

Overall, I really enjoyed working with my teammates. Konstantinos (K) worked really hard on coding and developing our final model with the neural network, incorporating all the relevant code (staying up till 7 a.m. one day), as well as fine tuning our parameters at the end to get us to our highest Kaggle score. Zach did significant work in getting the neural network up and running, and helped me understand the basics of it as well. Andrew helped out a lot in optimizing our pre-processing models (stop words, punctuation, lemmatization, etc.). I experimented with some of the preprocessing as well, playing around with the TF-IDF model (which ended up being surpassed by Word2Vec) for accumulating all the information regarding the text descriptions. I also tried running Word2Vec with linear regression. Regarding the report, I wrote about the two different models (Model 1 and Model 2), as well as proofread the entire report at the end.